

Machine Learning - Exercise 2

Robin Perälä 175910, Alexander Norrgran 164620, Robert Tommola 194000

12.4.2022

1.

ISLR, Exercise 5.4.2, p. 219-220. *We will now derive the probability that a given observation is part of a bootstrap sample. Suppose that we obtain a bootstrap sample from a set of n observations.*

- (a) *What is the probability that the first bootstrap observation is not the j th observation from the original sample? Justify your answer.*

If we have n observations and draw a random draw from them, the probability to obtain the j th observation (a specific one) is $1/n$. The complement of that is the probability to not obtain the j th observation. The complement is defined as $1 - 1/n = (n - 1)/n$.

Give formulas and use correct wording (permutations)...

- (b) *What is the probability that the second bootstrap observation is not the j th observation from the original sample?*

Bootstrap samples are drawn with replacement, so the probabilities of each draw is the same. The probability is $(n-1)/n$

- (c) *Argue that the probability that the j th observation is not in the bootstrap sample is $(1 - 1/n)^n$.*

If the probability that the j th observation is $(n-1)/n$ for one specific draw, then we compute the probability that is not in any of the draws by multiplying. $(n - 1)/n * (n - 1)/n * \dots * (n - 1)/n$ (not in the first draw, and not in the second draw, ... , an not in the last draw). Bootstrap sampling uses the same sample size as the original sample (n). This means that we have $((n-1)/n)^n$. $(n-1)/n$ can also be written as $1 - 1/n$, which means that we have $(1-1/n)^n$.

- (d) *When $n = 5$, what is the probability that the j th observation is in the bootstrap sample?*

We use the complement $1 - (1 - 1/n)^n = 1 - (1 - 1/5)^5 = 1 - (4/5)^5 = 0.67232$

- (e) *When $n = 100$, what is the probability that the j th observation is in the bootstrap sample?*


$1 - (99/100)^{100} = 0.6339677$

- (f) *When $n = 10,000$, what is the probability that the j th observation is in the bootstrap sample?*

$1 - (9999/10000)^{10000} = 0.632139$ When n grows bigger, the probability is getting ever closer to $1 - 1/e$

- (g) Create a plot that displays, for each integer value of n from 1 to 100,000, the probability that the j th observation is in the bootstrap sample. Comment on what you observe.

```
n = 1:100000
probs = 1 - ((n-1)/n)^{n}
plot(n, probs, log='xy')
```



- (h) We will now investigate numerically the probability that a bootstrap sample of size $n = 100$ contains the j th observation. Here $j = 4$. We repeatedly create bootstrap samples, and each time we record whether or not the fourth observation is contained in the bootstrap sample. Comment on the results obtained.

```
set.seed(1)
store=rep(NA, 10000)
for(i in 1:10000){
  store[i]=sum(sample(1:100, rep=TRUE)==4)>0
}
mean(store) #0.6405
```

```
## [1] 0.6417
```

```
#Which is getting ever closer to 1 - 1/e  
1-1/exp(1) #0.63212...
```

```
## [1] 0.6321206
```

2.

Suppose that $n = 10$ and the observations are 6.45, 1.28, -3.48, 2.44, -5.17, -1.67, -2.03, 3.58, 0.74, -2.14. Write a script in R to simulate the fraction of the original observations not contained in a bootstrap sample. Use $B = 10000$ bootstrap replications. Compare with the approximation $10/3$.

```
set.seed(1)
n=10
obs = c(6.45, 1.28, -3.48,
        2.44, -5.17, -1.67,
        -2.03, 3.58, 0.74, -2.14)

store=rep(NA, 10000)
for(i in 1:10000){
  store[i]= (n - sum(obs %in% sample(obs, size=10, rep=TRUE)))/n
}
mean(store)
```

```
## [1] 0.35057
```

```
#If we compare the mean value of observations included in the bootstrap (instead of the fraction) it is
mean(store)*n
```

```
## [1] 3.5057
```

```
#Which is close to
10/3
```

```
## [1] 3.333333
```

3

ISLR, Exercise 8.4.2, p. 361 *It is mentioned in Section 8.2.3 that boosting using depth-one trees (or stumps) leads to an additive model: that is, a model of the form*

$$f(X) = \sum_{j=1}^p f_j(X_j).$$

Explain why this is the case. You can begin with (8.12) in Algorithm 8.2.

Equation 8.12:

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x)$$

Because the tree contains just one split (stump), and as such the shrinkage parameter is 1? / 0?

Because the trees are only stumps, when we update \hat{f}^b by adding in a shrunken version of the new tree

$$\hat{f}^b(x) = \hat{f}^b(x) + \lambda \hat{f}^b(x)$$

, it only results in...

4.

Use the data `USCompaniesdata.dta`. Create a training set containing half of the observations, and a test set containing the remaining observations. Fit a tree with Return on Assets (`roa_w`) as the response and the other variables as predictors.

```
library(haven) #For importing Stata data
library(tree) #For fitting trees
library(randomForest) #For bagging and randomforests
library(gbm) #For boosting

#Load Stata data with read_dta function from haven
USCompanies_data_winsorized = read_dta("USCompanies_data_winsorized.dta")
USData = subset(USCompanies_data_winsorized, select = -conm)
USData = na.omit(USData)

#Split into training and test set
set.seed(1)
train <- sample(1:nrow(USData), nrow(USData)/2)
test <- (-train)

# set.seed(1)
# train=sample(c(TRUE, FALSE), nrow(USData), rep=TRUE)
# test = !train
```

```
#Fit a tree using function "tree" (in library "tree")
treeROA = tree(roa_w ~ ., USData, subset=train)
summary(treeROA)
```

```
##
## Regression tree:
## tree(formula = roa_w ~ ., data = USData, subset = train)
## Variables actually used in tree construction:
## [1] "profit_margin_w" "icapt_w" "at_w" "ebitda_w"
## [5] "roe_w"
## Number of terminal nodes: 12
## Residual mean deviance: 0.01093 = 18.31 / 1676
## Distribution of residuals:
##      Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## -0.6040000 -0.0382000  0.0001005  0.0000000  0.0439200  0.8494000
```


```
treeROA
```

```
## node), split, n, deviance, yval
##      * denotes terminal node
##
##  1) root 1688 121.2000 -0.085170
##    2) profit_margin_w < -0.910934 311 33.0700 -0.486700
##      4) icapt_w < 12.195 102 8.1060 -0.763400
##        8) at_w < 4.345 43 2.1370 -0.903900 *
##        9) at_w > 4.345 59 4.5000 -0.660900
##       18) ebitda_w < -3.864 42 1.6620 -0.788500 *
```

```
##      19) ebitda_w > -3.864 17    0.4666 -0.345800 *
##      5)  icapt_w > 12.195 209   13.3500 -0.351700
##      10) roe_w < -0.452831 96    5.3960 -0.505900
##      20) profit_margin_w < -3.15281 52    2.4670 -0.611800 *
##      21) profit_margin_w > -3.15281 44    1.6560 -0.380700 *
##      11) roe_w > -0.452831 113    3.7360 -0.220800
##      22) roe_w < 0.67636 98    0.9485 -0.172600 *
##      23) roe_w > 0.67636 15    1.0750 -0.535500 *
##      3) profit_margin_w > -0.910934 1377    26.6800 0.005529
##      6) ebitda_w < -0.0815 319    9.5560 -0.174000
##      12) at_w < 11.494 39    2.9000 -0.422400 *
##      13) at_w > 11.494 280    3.9150 -0.139400
##      26) roe_w < -0.430403 72    1.5760 -0.251500 *
##      27) roe_w > -0.430403 208    1.1220 -0.100600 *
##      7) ebitda_w > -0.0815 1058    3.7430 0.059660
##      14) roe_w < 0.118579 534    1.1950 0.023110 *
##      15) roe_w > 0.118579 524    1.1070 0.096910 *
```


#Plot

```
plot(treeROA); text(treeROA, pretty=0, cex=0.7, srt=25)
```



#MSE

```
yhat=predict(treeROA, newdata=USData[-train,])
ROAtest=USData[-train, "roa_w"]
plot(yhat, ROAtest$roa_w, ylab = "roa_w")
abline(0,1)
```



```
mean((yhat-ROAtest$roa_w)^2)
```

```
## [1] 0.01246236
```


5.

Apply bagging to *USCompaniesdata.dta*. Compare the MSE of the tree in Exercise 4 with the MSE of the bagged trees.

#Bagging

```
#Bagging using the function randomForest,
#inside the randomForest library. When mtry is for all variables, it is bagging.
bagROA = randomForest(roa_w~., USData, subset=train, mtry=ncol(USData)-1, importance =TRUE)
summary(bagROA)
```


```
##               Length Class  Mode
## call           6      -none- call
## type           1      -none- character
## predicted      1688     -none- numeric
## mse            500     -none- numeric
## rsq            500     -none- numeric
## oob.times      1688     -none- numeric
## importance      88      -none- numeric
## importanceSD    44      -none- numeric
## localImportance 0       -none- NULL
## proximity       0       -none- NULL
## ntree           1       -none- numeric
## mtry            1       -none- numeric
## forest          11      -none- list
## coefs           0       -none- NULL
## y              1688     -none- numeric
## test            0       -none- NULL
## inbag           0       -none- NULL
## terms           3       terms  call
```

bagROA

```
##
## Call:
## randomForest(formula = roa_w ~ ., data = USData, mtry = ncol(USData) -      1, importance = TRUE, s
##               Type of random forest: regression
##               Number of trees: 500
## No. of variables tried at each split: 44
##
##               Mean of squared residuals: 0.007524909
##               % Var explained: 89.52
```

#MSE

```
yhat=predict(bagROA, newdata=USData[-train,])
ROAtest=USData[-train, "roa_w"]
plot(yhat, ROAtest$roa_w)
abline(0,1)
```



```
mean((yhat-ROAtest$roa_w)^2)
```

```
## [1] 0.006178823
```

6.

Apply random forests to *USCompaniesdata.dta*. Does random forests provide an improvement over the bagged trees in Exercise 5?

```
#Random forests
```

```
set.seed(1)
#Fit a random forest using the function randomForest,
#inside the randomForest library. When mtry is less than
#than all variables in the data, it is a random forests model
forestROA = randomForest(roa_w~., USData, subset=train, importance = TRUE)
# forestROA = randomForest(roa_w~., USData, subset=train, mtry=ncol(USData)/2, importance = TRUE)
# forestROA = randomForest(roa_w~., USData, subset=train, mtry=ncol(USData)/3, importance = TRUE)
#Should be done on several different mtry

summary(forestROA)
```

```
##              Length Class  Mode
## call              5  -none- call
## type              1  -none- character
## predicted        1688  -none- numeric
## mse              500  -none- numeric
## rsq              500  -none- numeric
## oob.times        1688  -none- numeric
## importance         88  -none- numeric
## importanceSD       44  -none- numeric
## localImportance    0  -none- NULL
## proximity          0  -none- NULL
## ntree              1  -none- numeric
## mtry               1  -none- numeric
## forest            11  -none- list
## coefs              0  -none- NULL
## y                 1688  -none- numeric
## test              0  -none- NULL
## inbag              0  -none- NULL
## terms              3   terms  call
```


```
forestROA
```

```
##
## Call:
## randomForest(formula = roa_w ~ ., data = USData, importance = TRUE,      subset = train)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 14
##
##              Mean of squared residuals: 0.007853636
##              % Var explained: 89.06
```

```
#MSE
```

```
yhat=predict(forestROA, newdata=USData[-train,])
```

```
ROAtest=USData[-train, "roa_w"]  
plot(yhat, ROAtest$roa_w)  
abline(0,1)
```



```
mean((yhat-ROAtest$roa_w)^2)
```


```
## [1] 0.006525587
```

7.

Apply boosting to *USCompaniesdata.dta*. Which variables are the most important predictors in the boosted model?

```
#Boosting
set.seed(1)
#Boosting using the "gbm" function inside the "gbm" library
boostROA = gbm(roa_w~., USData, distribution="gaussian", n.trees=500, interaction.depth=4) #5000

summary(boostROA)
```



```
##                                var      rel.inf
## profit_margin_w      profit_margin_w 49.077220066
## ebitda_w            ebitda_w 12.808796933
## roe_w              roe_w 8.370929651
## icapt_w            icapt_w 7.328434141
## at_w              at_w 4.762816673
## teq_w            teq_w 4.598167229
## ib_w            ib_w 1.470522692
## asset_turnover_w  asset_turnover_w 1.463603848
## lt_w            lt_w 1.194587644
## pe_w            pe_w 0.863484862
## ch_w            ch_w 0.656509777
```


```
## sale_w          sale_w 0.655561154
## oancf_w         oancf_w 0.613927694
## cogs_w          cogs_w 0.595359456
## operating_margin_w operating_margin_w 0.555910076
## bkvlp_w         bkvlp_w 0.495882428
## gdwl_w          gdwl_w 0.467283587
## ap_w            ap_w 0.352596398
## dp_w            dp_w 0.312692577
## dlc_w           dlc_w 0.296799663
## fopo_w          fopo_w 0.272657843
## lco_w           lco_w 0.247277590
## txt_w           txt_w 0.243039866
## ppent_w         ppent_w 0.219257314
## ceq_w           ceq_w 0.217857303
## aco_w           aco_w 0.190576224
## chech_w         chech_w 0.182728402
## ci_w            ci_w 0.174844669
## rect_w          rect_w 0.158690123
## intano_w        intano_w 0.147219421
## caps_w          caps_w 0.142781679
## ivncf_w         ivncf_w 0.141636410
## capx_w          capx_w 0.139970966
## re_w            re_w 0.125825184
## np_w            np_w 0.120447364
## epspi_w         epspi_w 0.119101764
## invt_w          invt_w 0.113414751
## fiaow_w         fiaow_w 0.046955740
## tstkw_w         tstkw_w 0.022022188
## dvt_w           dvt_w 0.016238275
## ivst_w          ivst_w 0.006382126
## aqc_w           aqc_w 0.005583599
## siv_w           siv_w 0.004404648
## cidergl_w       cidergl_w 0.000000000
```

```
boostROA
```

```
## gbm(formula = roa_w ~ ., distribution = "gaussian", data = USData,
##       n.trees = 500, interaction.depth = 4)
## A gradient boosted model with gaussian loss function.
## 500 iterations were performed.
## There were 44 predictors of which 43 had non-zero influence.
```

```
#all or only training set?
```

```
#MSE
yhat=predict(boostROA, newdata=USData[-train,])
ROAtest=USData[-train, "roa_w"]
plot(yhat, ROAtest$roa_w)
abline(0,1)
```



```
mean((yhat-ROAtest$roa_w)^2)
```

```
## [1] 0.0008299824
```