

Machine Learning - Exercise 4: Second part

Robin Perälä 175910

5.5.2022

Task 5

Adapted from R-bloggers (2021): Deep Neural Network in R. r-bloggers.com/2021/04/deep-neural. Fit a neural network with a single hidden layer to the Boston Housing data.

```
rm(list=ls()); RNGkind(sample.kind='Rejection')

#Load Library
library(mlbench) #Contains BostonHousing data
library(dplyr) #Contains "mutate_if" function
library(magrittr) #Contains pipes, e.g. %<>%
library(neuralnet) #For fitting neural networks
library(keras) #Contains "keras_model_sequential" function
library(glmnet) #For regularized linear models (e.g. lasso and ridge)

#Getting Data
data("BostonHousing")
data <- BostonHousing
str(data)

## 'data.frame': 506 obs. of 14 variables:
## $ crim : num 0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn : num 18 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus : num 2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ nox : num 0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm : num 6.58 6.42 7.18 7 7.15 ...
## $ age : num 65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis : num 4.09 4.97 4.97 6.06 6.06 ...
## $ rad : num 1 2 2 3 3 3 5 5 5 ...
## $ tax : num 296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ b : num 397 397 393 395 397 ...
## $ lstat : num 4.98 9.14 4.03 2.94 5.33 ...
## $ medv : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

```
data %<>% mutate_if(is.factor, as.numeric)
data <- as.matrix(data)
```

#Data Partition

```

set.seed(707)
ind <- sample(2, nrow(data), replace = T, prob = c(.7, .3))
training <- data[ind==1, 1:13]
test <- data[ind==2, 1:13]
trainingtarget <- data[ind==1, 14]
testtarget <- data[ind==2, 14]

#Scaling
m <- colMeans(training)
s <- apply(training, 2, sd)
training <- scale(training, center = m, scale = s)
test <- scale(test, center = m, scale = s)

#Model Creation
model <- keras_model_sequential()
model %>%
  layer_dense(units = 5, activation = 'relu', input_shape = c(13)) %>%
  layer_dense(units = 1)

#Model Compilation
model %>% compile(loss = 'mse',
                  optimizer = 'rmsprop',
                  metrics = 'mae')

#Model Fitting
mymodel <- model %>%
  fit(training, trainingtarget,
      epochs = 200,
      batch_size = 32,
      validation_split = 0.2)

#Prediction
model %>% evaluate(test, testtarget)

##      loss      mae
## 42.972755  4.620593


pred <- model %>% predict(test)

#MSE and MAE (The same ones as above)
mean((testtarget-pred)^2); mean(abs(testtarget-pred))


## [1] 42.97276
## [1] 4.620593

#Scatter Plot Original vs Predicted
plot(testtarget, pred)

```



```
#Plot model error for each epoch  
plot(mymodel)
```



Task 6

What does `activation <- nn_relu()` mean? What does `dropout <- nn_dropout(0.4)` mean?

nn_relu() is specifying the rectified linear unit function as our activation function of choice. This activation function is then used in our neural network.

ReLU activation function is defined as

$$g(z) = \begin{cases} z, & z > 0 \\ 0, & z \leq 0 \end{cases}$$

nn_dropout(0.4) is randomly zeroing some of the units with a probability of 40%. This regularization can be applied to avoid overfitting

Task 7

Experiment with adding a second hidden layer. Does it improve the test performance?

```
#Model Creation
model2 <- keras_model_sequential()
model2 %>%
  layer_dense(units = 5, activation = 'relu', input_shape = c(13)) %>%
  layer_dropout(rate=0.3) %>%
  layer_dense(units = 10, activation = 'relu') %>%
  layer_dropout(rate=0.2) %>%
  layer_dense(units = 1)

#Model Compilation
model2 %>% compile(loss = 'mse',
                  optimizer = 'rmsprop',
                  metrics = 'mae')

#Model Fitting
mymodel2 <- model2 %>%
  fit(training, trainingtarget,
      epochs = 200,
      batch_size = 32,
      validation_split = 0.2)

#Prediction
model2 %>% evaluate(test, testtarget)
```

```
##      loss      mae
## 29.806849  3.927519
```


```
pred <- model2 %>% predict(test)

#MSE and MAE (The same ones as above)
mean((testtarget-pred)^2); mean(abs(testtarget-pred))
```

```
## [1] 29.80685
```


```
## [1] 3.927519
```

```
#Scatter plot Original vs Predicted
plot(testtarget, pred)
```



```
#Plot model error for each epoch  
plot(mymodel2)
```

```
## 'geom_smooth()' using formula 'y ~ x'
```



A second layer with 10 nodes seems to improve prediction performance. However, I also tested a model with 5 nodes in the second layer, which seemed to perform worse than the model with a single layer.

Task 8

Select a final neural network. Compare the MAE and MSE on the test set with linear regression with regularisation.

```
#Neural network with one hidden layer
```

```
model %>% evaluate(test, testtarget)
```

```
##      loss      mae
## 42.972755  4.620593
```

```
pred <- model %>% predict(test)
```

```
#MSE and MAE (The same ones as above)
```

```
NN1MSE = mean((testtarget-pred)^2)
```

```
NN1MAE = mean(abs(testtarget-pred))
```

```
#Neural network with two hidden layers
```

```
model2 %>% evaluate(test, testtarget)
```

```
##      loss      mae
## 29.806849  3.927519
```

```
pred <- model2 %>% predict(test)
```

```
#MSE and MAE (The same ones as above)
```

```
NN2MSE = mean((testtarget-pred)^2)
```

```
NN2MAE = mean(abs(testtarget-pred))
```

```
#Use function glmnet (from library glmnet) to perform linear regression with Lasso regularisation  
#Alpha=1 means lasso.
```

```
#We use it on the same training data as was used for the neural network models
```

```
lassoModel = glmnet(x=training, y=trainingtarget, alpha=1, standardize=TRUE)
```

```
#Use 10-fold cross-validation (on the training data) for finding optimal lambda
```

```
lassoCV=cv.glmnet(x=training, y=trainingtarget, alpha=1, standardize=TRUE, nfolds=10)
```

```
#Test lasso on the same test data as used for the neural network models
```

```
lassoPred=predict(lassoModel, s=lassoCV$lambda.min, newx=test)
```

```
#MSE and MAE
```

```
lassoMSE = mean((lassoPred-testtarget)^2)
```

```
lassoMAE = mean(abs(lassoPred-testtarget))
```

```
#For nicer output (printing)
```

```
comparisonMatrix = matrix(c(NN1MSE, NN1MAE,  
                             NN2MSE, NN2MAE,  
                             lassoMSE, lassoMAE),  
                           ncol=2,
```

```

                                byrow=TRUE)
comparisonMatrix = round(comparisonMatrix, 4)

colnames(comparisonMatrix) = c("MSE", "MAE")
rownames(comparisonMatrix) = c("Neural network with one hidden layer",
                                "Neural network with two hidden layers",
                                "Linear regression with Lasso regularisation")

print(comparisonMatrix)

```

```

##                                MSE    MAE
## Neural network with one hidden layer    42.9728 4.6206
## Neural network with two hidden layers    29.8069 3.9275
## Linear regression with Lasso regularisation 21.9208 3.5699

```

Lasso seems to have the best prediction performance. The performance of the neural network might improve if different parameters are used E.g. I noticed that using a lower dropout rate and more neurons in the second layer improved performance. The number of epochs for training is another factor.